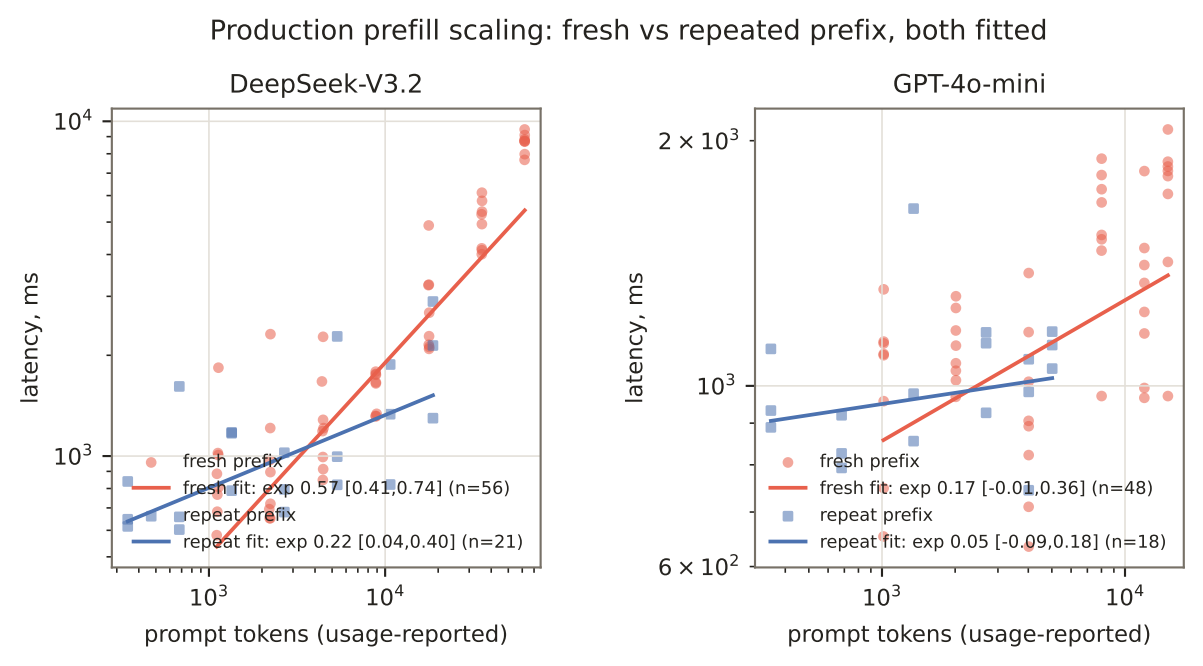
Production prefill scaling: fresh vs repeated prefix, both fitted
DeepSeek-V3.2
GPT-4o-mini
latency, ms
prompt tokens (usage-reported)
fresh prefix
fresh fit: exp 0.57 [0.41,0.74] (n=56)
repeat prefix
repeat fit: exp 0.22 [0.04,0.40] (n=21)
fresh fit: exp 0.17 [-0.01,0.36] (n=48)
repeat fit: exp 0.05 [-0.09,0.18] (n=18)